



Q learning



Roger Garcia
Mississippi State University



Reinforcement Learning (Model-Free)

- Off-Policy vs On-Policy

- Function that maps states to actions
- Probability of an action given a state

$$\pi: S \rightarrow A$$

- MDP?

- A set of states
- A transition model $Q(\text{state}, \text{action}) = R(\text{state}, \text{action}) + \gamma \max_{a'} [Q(\text{next_s}, a')]$
- A reward function
- Acquire optimal policy [after convergence]
- Off-Policy: Must actually try actions and states out to learn !

$$a = \pi(s)$$

- Value iteration

- Utility value [Q value] = $Q(s, a)$ for every action at every state
- Utility of state s ?
 - maxQ value over all possible actions at that state
- Utilities of actions?
 - Progressively improve utilities of actions

$$a_t = \pi(s_t) = \operatorname{argmax}_a Q(s_t, a)$$

Q learning

- Q learning implementation using Q-Table
 - Specifies the value of each action for each possible state
 - Helps map from values received from the environment to Q-Table indices.
- There are 4 possible actions: [up, down, left, right]
- Objective: To move from starting position to the green goal position
- Search Policy
 - e-Greedy; the robot chooses its next action as the best action it knows at the time
 - Look at the Q-Value for each action at the current state
- **DEMO TIME** (10x10 GridWorld)(<https://github.com/aalind0/RL-Game-Bot>)
 - Act randomly in search of reward
 - Remember the reward we get for being in state and taking Act randomly in search of a reward
 - The reward for the previous state and action is a combination of the current immediate reward of moving forward and also the best possible reward of any action from the current state, the future reward.
 - Prefer paths which result in positive rewards and avoid negative ones

The Q-Learning Algorithm

- Parameters
- Initialize
- For each World.iteration:
 - Select rand initial state
 - While not at Goal:
 - Select 1 among all actions at curr_state
 - Valid action? Consider transitioning to next state
 - Get maxQ value for this next state based on all possible actions
 - Compute $Q(s, action)$ and update
 - Set the next state as the curr_state
 - End
- End

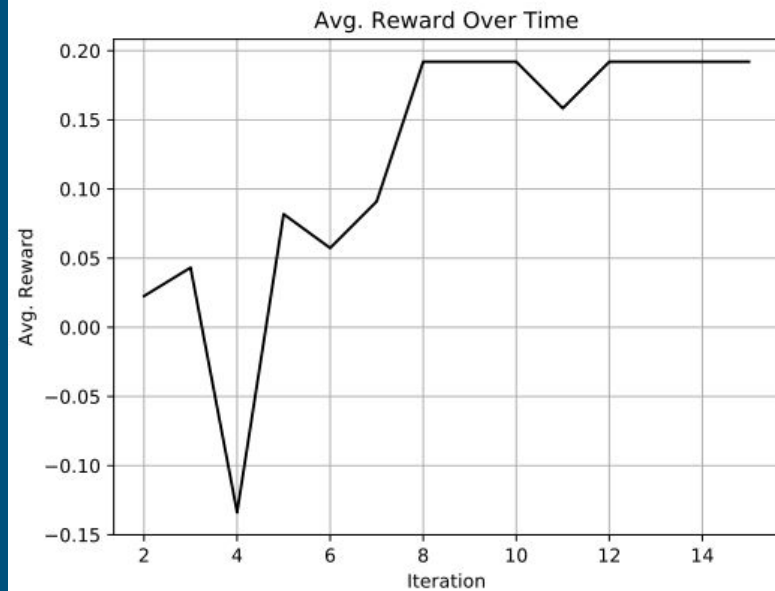
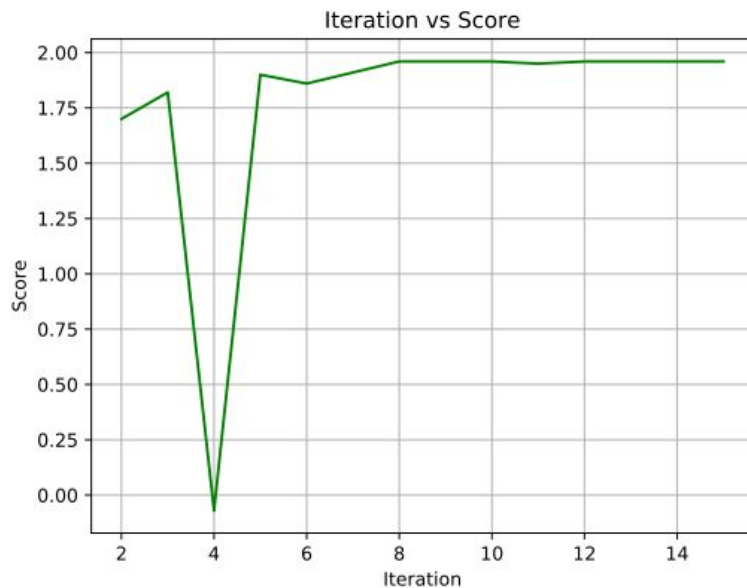
**Learning Rate
Vs
Discount Factor**

Q Learning Simulation Runs (Grid World)



- Discount Factor (Gamma): **0.3**
- Initial Score, alpha : **1.0**
- Walk Reward : - **0.01**
- Red Square (you die, you restart) : -**1.0**
- Goal Square : + **1.0**
- Actions: LEFT, DOWN, RIGHT, UP
- Policy ?
 - Explore environment ? random actions
 - Use Greedy Policy to evaluate states
 - Update gather information, [progressively reduce randomness]
- Highest Reward = Action to take
 - Take random action uniformly
 - If enough World iterations are completed, each action will be tried an infinite number of times, thus ensuring optimal actions to be discovered

**Goal:
Maximize
Sum of Rewards**



5x5 Grid World



5x5 Grid World




10x10 Grid World




10x10 Grid World


Avg. Reward Over Time



Avg. α Over Time

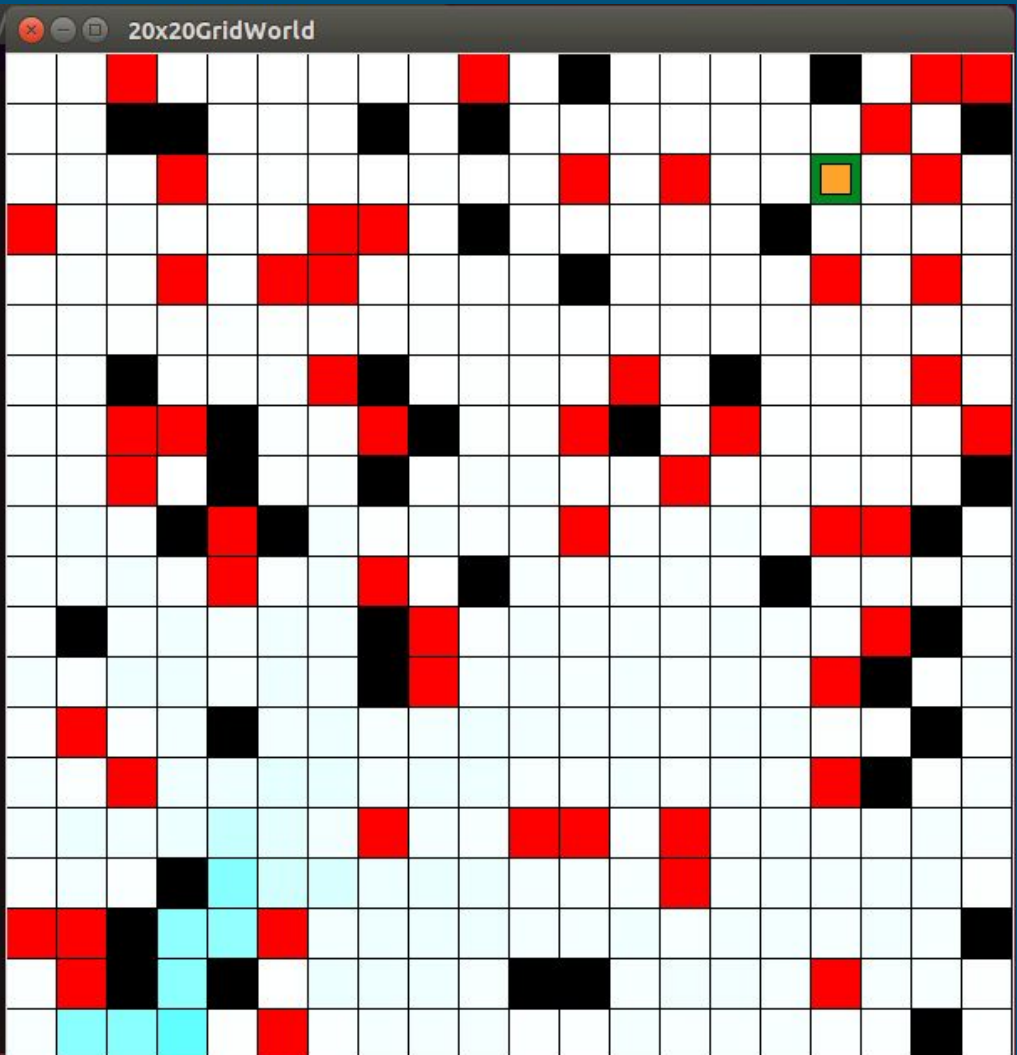


Decay α over Time (optimal)




20x20 Grid World

```
roger@garcroge: ~/Documents/spring2017/
| roger@garcroge: ~/Documents/spring2017/specialtopic_
| *****
| World Iter: 121
| Move: 150
| *****
| ACTION: right
| ALPHA: 0.605885939418
| SCORE: -0.5
| Status: [FAILURE] TRY AGAIN
| *****
| World Iter: 121
| Move: 151
| *****
| ACTION: down
| ALPHA: 0.605483489637
| SCORE: -0.51
| Status: [FAILURE] TRY AGAIN
| *****
| World Iter: 121
| Move: 152
| *****
| ACTION: right
| ALPHA: 0.60508396099
| SCORE: 0.49
| Status: [SUCCESS] GOAL
| *****
| Average Learning Rate: 0.670338349627
| Discount Factor (gamma): 0.3
| Walk Reward: -0.01
| World Iter: 122
| Number of Moves (actions): 152
| *****
| PAUSING
|
```




20x20 Grid World




20x20 Grid World

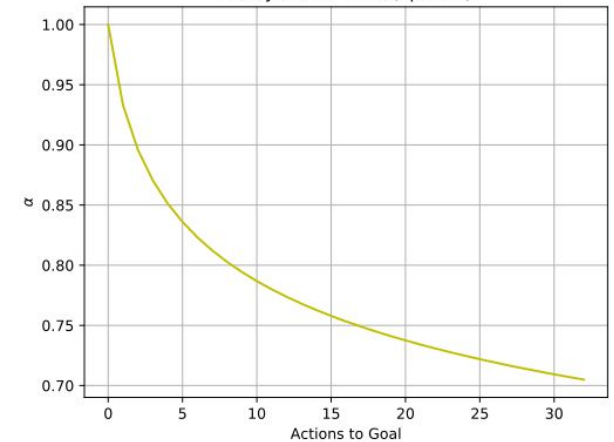
Avg. Reward Over Time




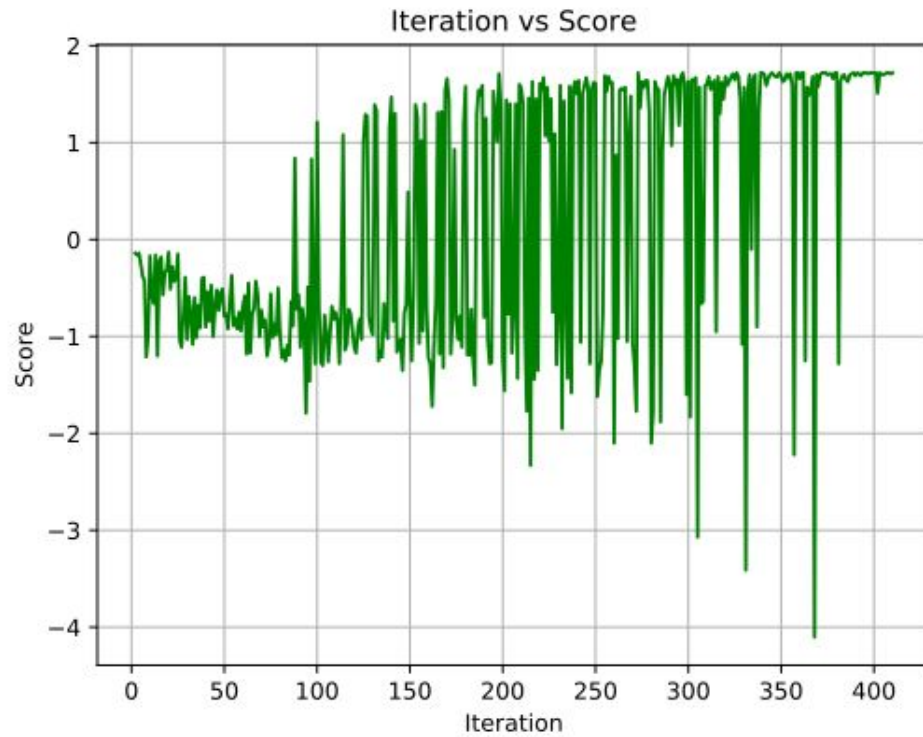
Avg. α Over Time




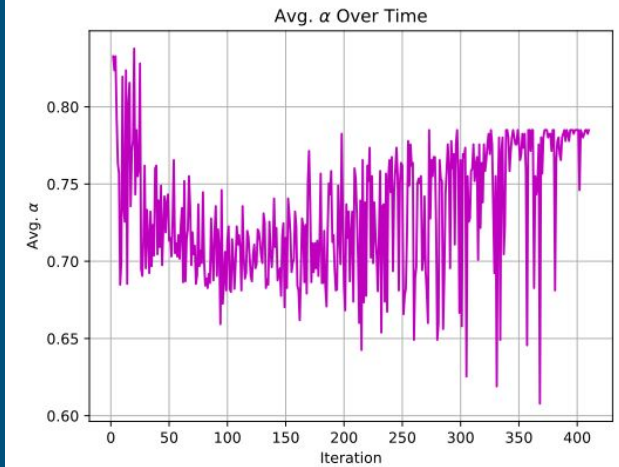

Decay α over Time (optimal)



30x30 Grid World

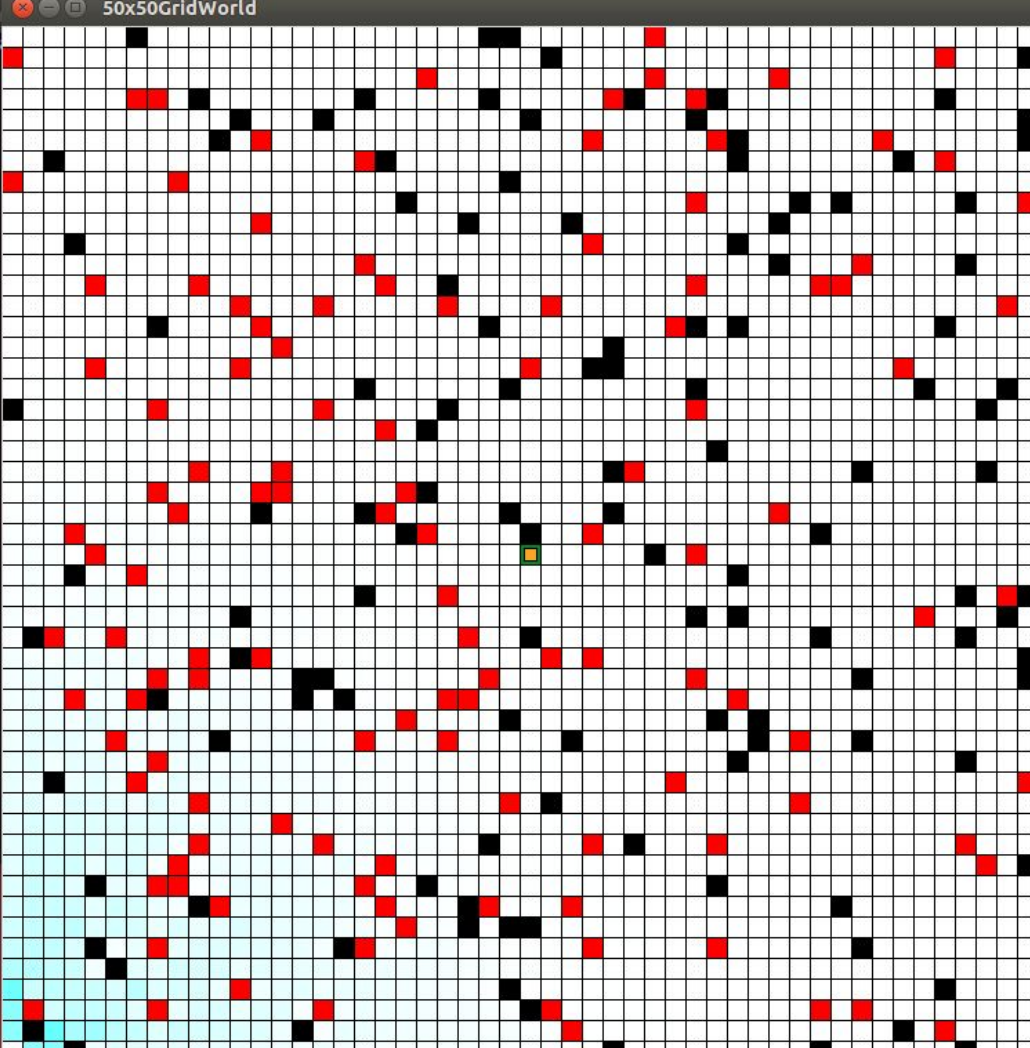


30x30 Grid World


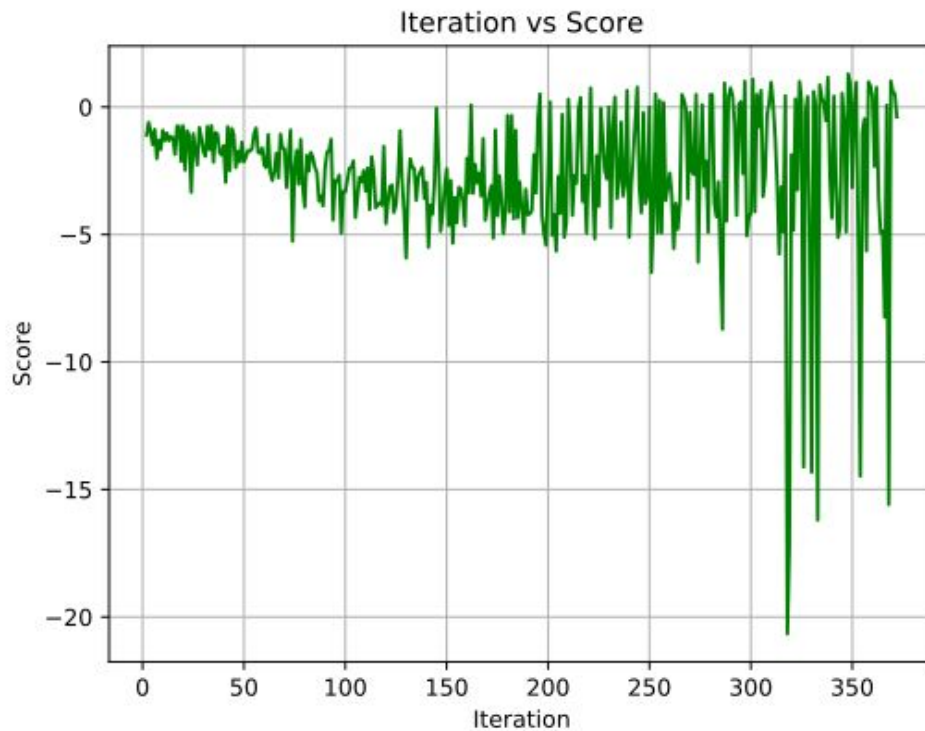


50x50 Grid World

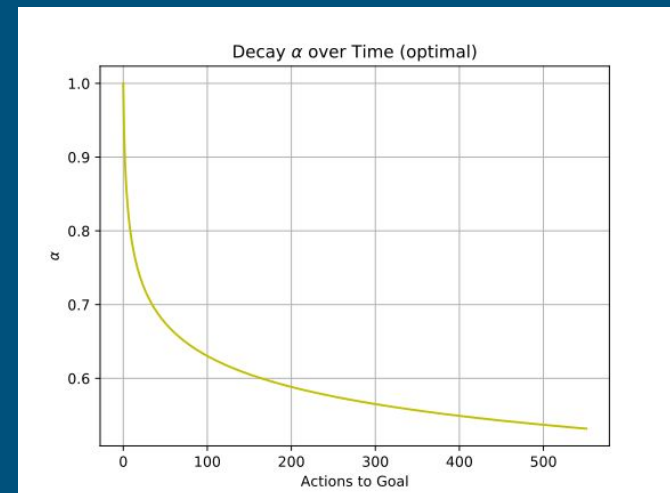
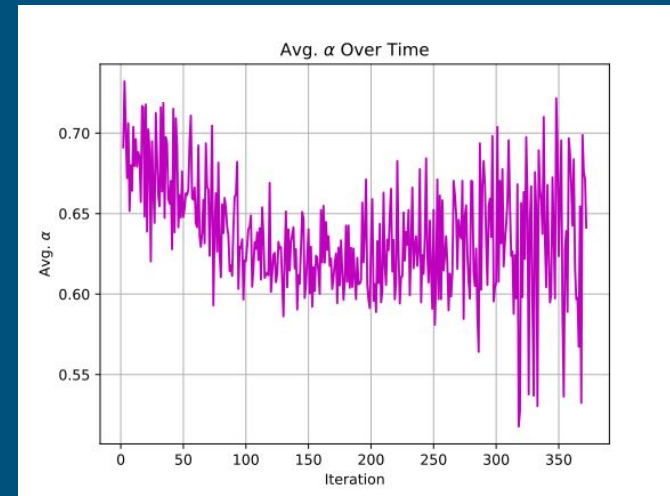
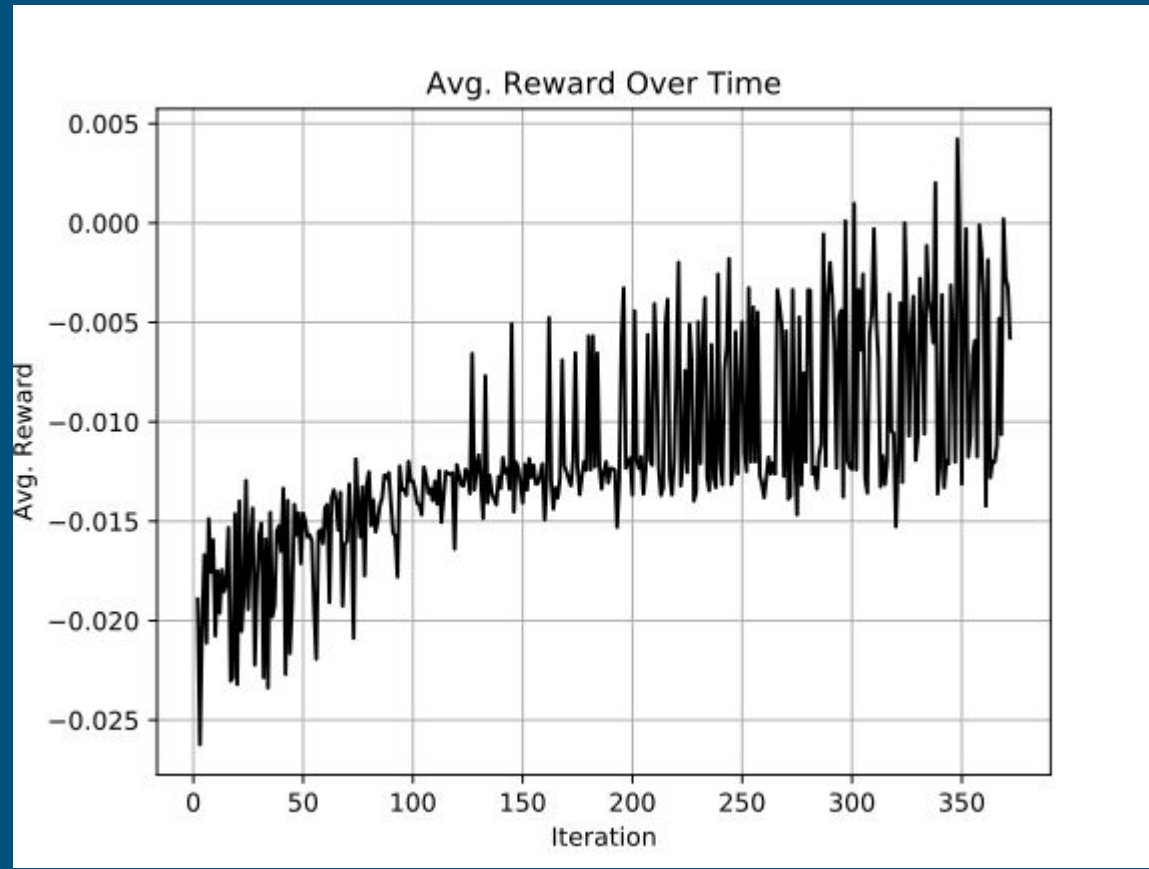
```
roger@garcroge: ~/Documents/spr... 50x50GridWorld
[roger@garcroge: ~/Documents/spring2017/speci
Status: [FAILURE] TRY AGAIN
*****
World Iter: 228
Move: 130
*****
ACTION: down
ALPHA: 0.614618551514
SCORE: -0.3
Status: [FAILURE] TRY AGAIN
*****
World Iter: 228
Move: 131
*****
ACTION: up
ALPHA: 0.614147757544
SCORE: -0.31
Status: [FAILURE] TRY AGAIN
*****
World Iter: 228
Move: 132
*****
ACTION: right
ALPHA: 0.613680900314
SCORE: 0.69
Status: [SUCCESS] GOAL
*****
Average Learning Rate: 0.679623456318
Discount Factor (gamma): 0.3
Walk Reward: -0.01
World Iter: 229
Number of Moves (actions): 132
*****
PAUSING
█
```



50x50 Grid World



50x50 Grid World



Conclusion

- Updating policy *always* even when exploring far from Goal
- Greedy Policy used (with exploration)
 - e-soft
 - Softmax (rank or weight)
- Future work
 - Explore tradeoff between policies, to produce best results overall
 - Impact of policy to learning
 - On-Policy vs Off-Policy
- References
 - <http://mnemstudio.org/path-finding-q-learning-tutorial.htm>
 - <https://github.com/aalind0/RL-Game-Bot>
 - <http://mbb-team.github.io/VBA-toolbox/wiki/Fast-demo-Q-learning-model/>
 - <http://www.cse.unsw.edu.au/~cs9417ml/RL1/tdlearning.html#aselection>